

# Agenda 6 October

Read Chapters 28 – 32 (Concurrency): Module 5

Quiz Review


Attendance

Synchronization Problems

Introduction to Locks

Test\_and\_set compare\_and\_swap

Quiz 2

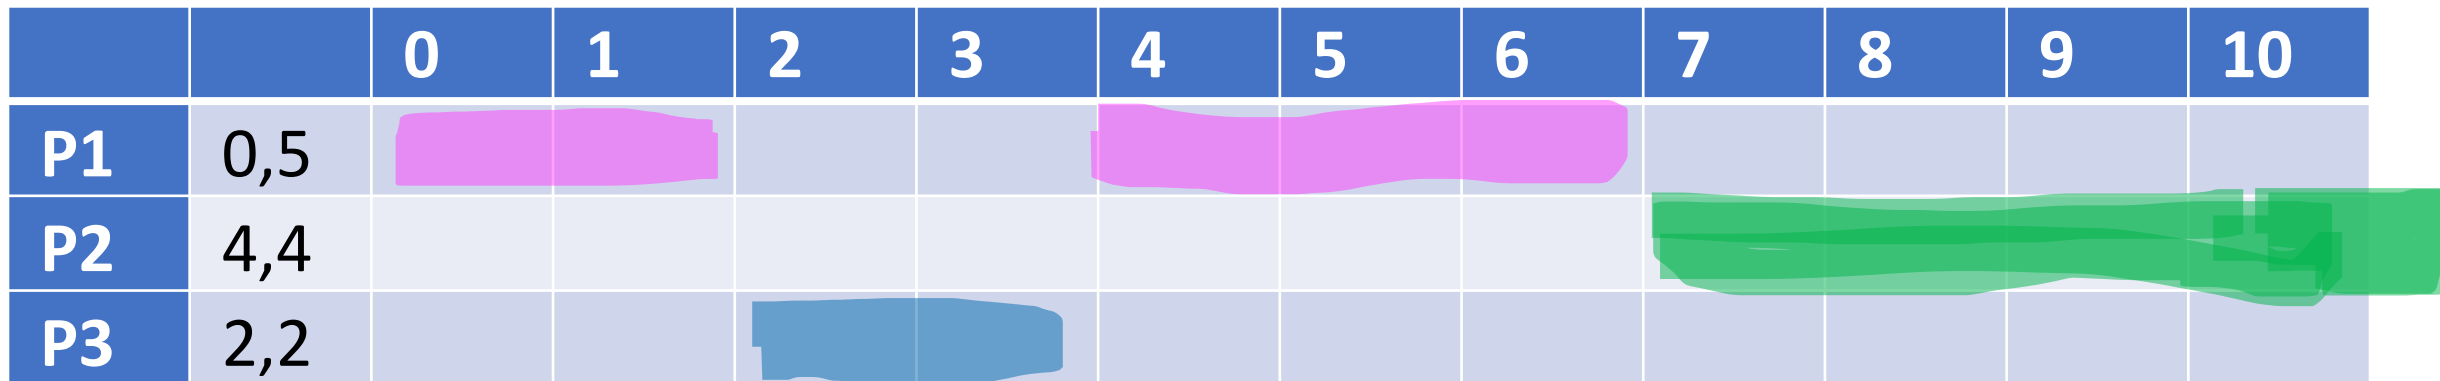


Three processes are executed using a **SCTF** scheduler:

Process	Arrival time	Total CPU time
P1	0	5
P2	4	4
P3	2	2

At what time does P1 end?

What is the SCTF algorithm or rules? Arrival time preemption  
& evaluate remaining process times



`fork()`: creates a new process, copying code and data to the child

As the processes continue to run, their memory spaces and variables are separate (independent)

Thread: all threads in the same process share all variables in the address space of the process.

`exec()` overlays the code and data of the process. Any code in the process following the `exec()` no longer exists.

When a `fork()` is given, 2 processes continue to run. If these processes `fork()` again, there are then 4 processes running.

```
fork(); statement a; statement b; fork(); statement c; statement d;
```

Real time in a system: time from initiation to completion: wait time + execution time

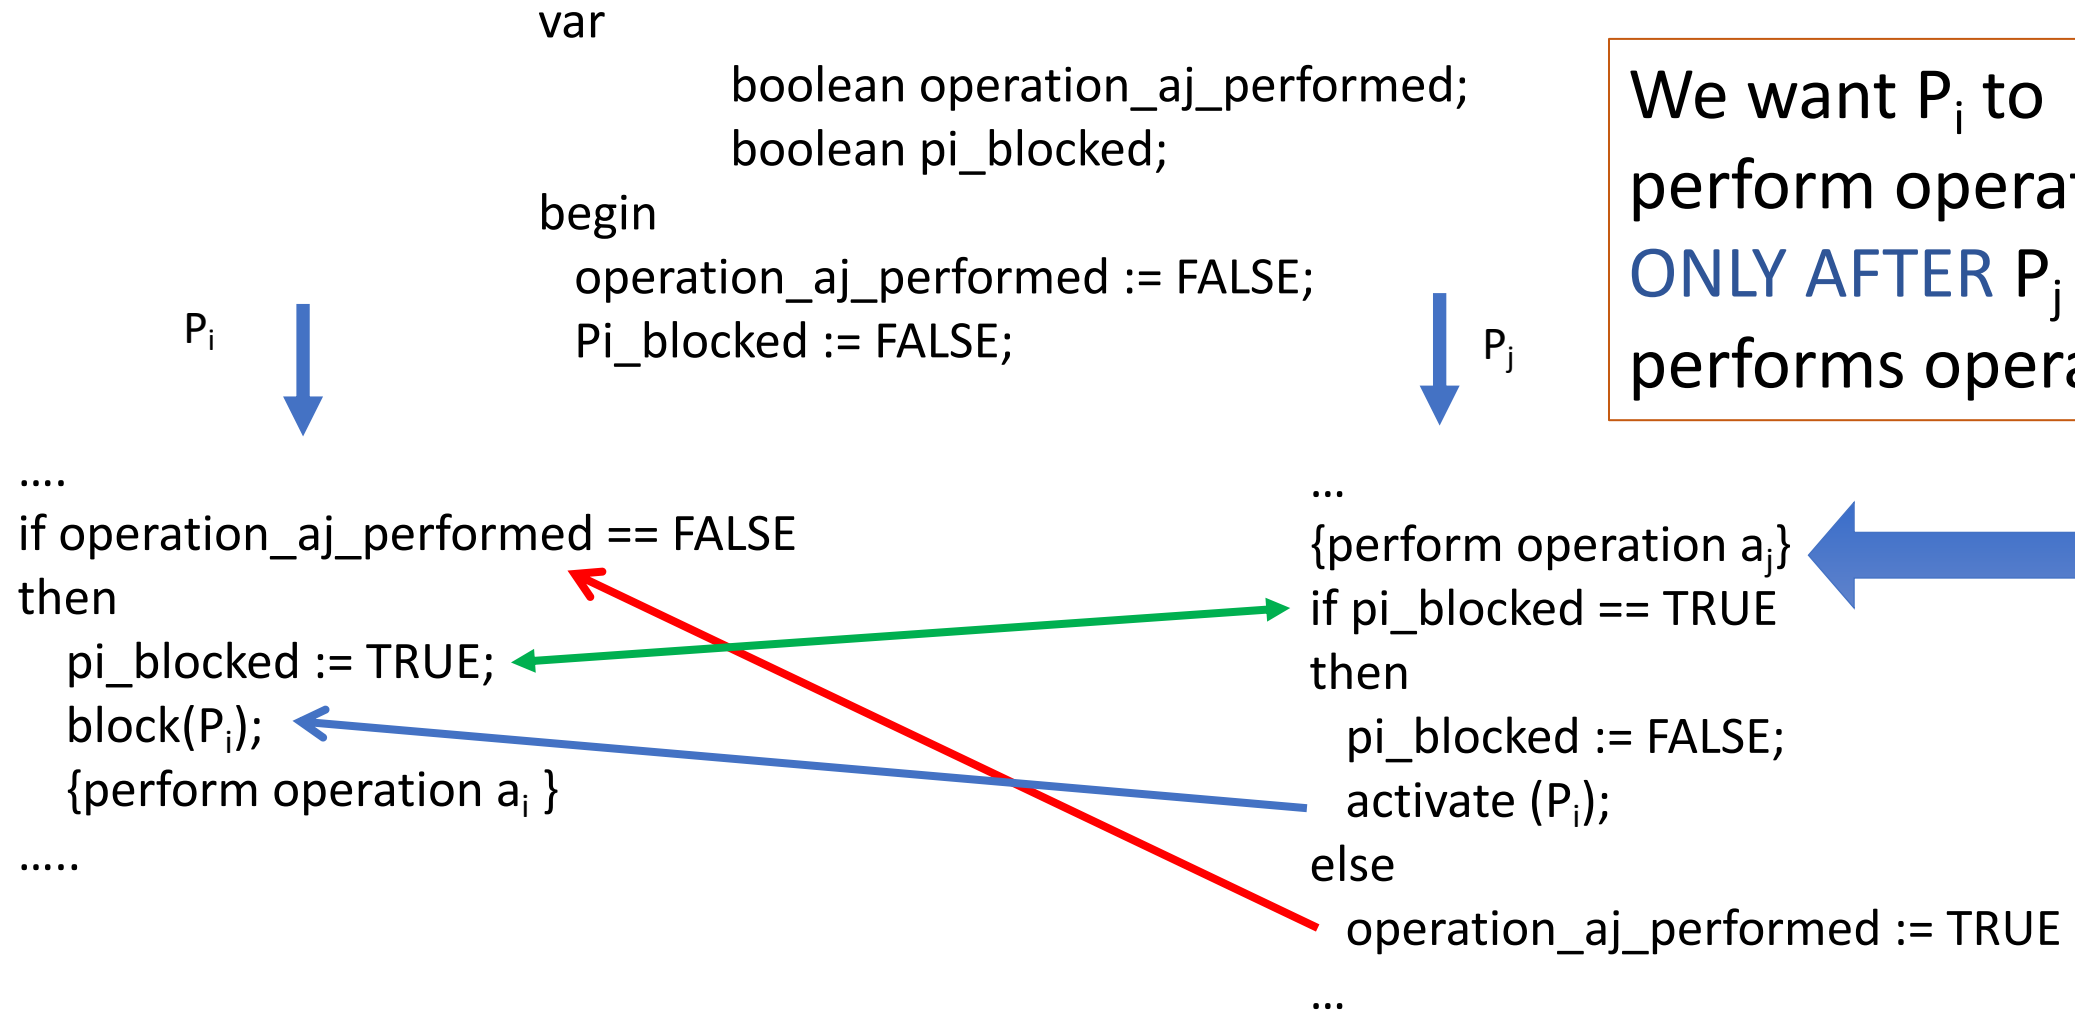
# Attempt at a 'Signaling' Solution for Synchronizing

Two Processes Cooperative in the solution of a problem

One generates data for the other to use.

Since they run asynchronously, we need a solution that is independent of the order of execution of their statements.

# A high level example Requiring a Coarse-grained Atomic Action



# The 'Signaling' Solution Fails

- $P_i$  may face indefinite blocking in some situation
- Use *indivisible or atomic operations* instead

Time	Actions of process $P_i$	Actions of process $P_j$
$t_1$	<b>if</b> $action_{aj\_performed} == false$	
$t_2$		{perform action $a_j$ }
$t_3$		<b>if</b> $pi\_blocked == true$
$t_4$		$action_{aj\_performed} := true$
$\vdots$		
$t_{20}$	$pi\_blocked := true;$	
$t_{21}$	$block(P_i);$	



# Update to the 'Signaling' Solution

- Atomic  $\Leftrightarrow$  Indivisible
- Two **atomic functions** make the solution work

Procedure check\_aj

begin

if operation\_aj\_performed = FALSE

then

pi\_blocked := TRUE;

block(P<sub>i</sub>);

end

Procedure post\_aj

begin

if pi\_blocked = TRUE

then

pi\_blocked := FALSE;

activate(P<sub>j</sub>)

else

operation\_aj\_performed := true;

end;

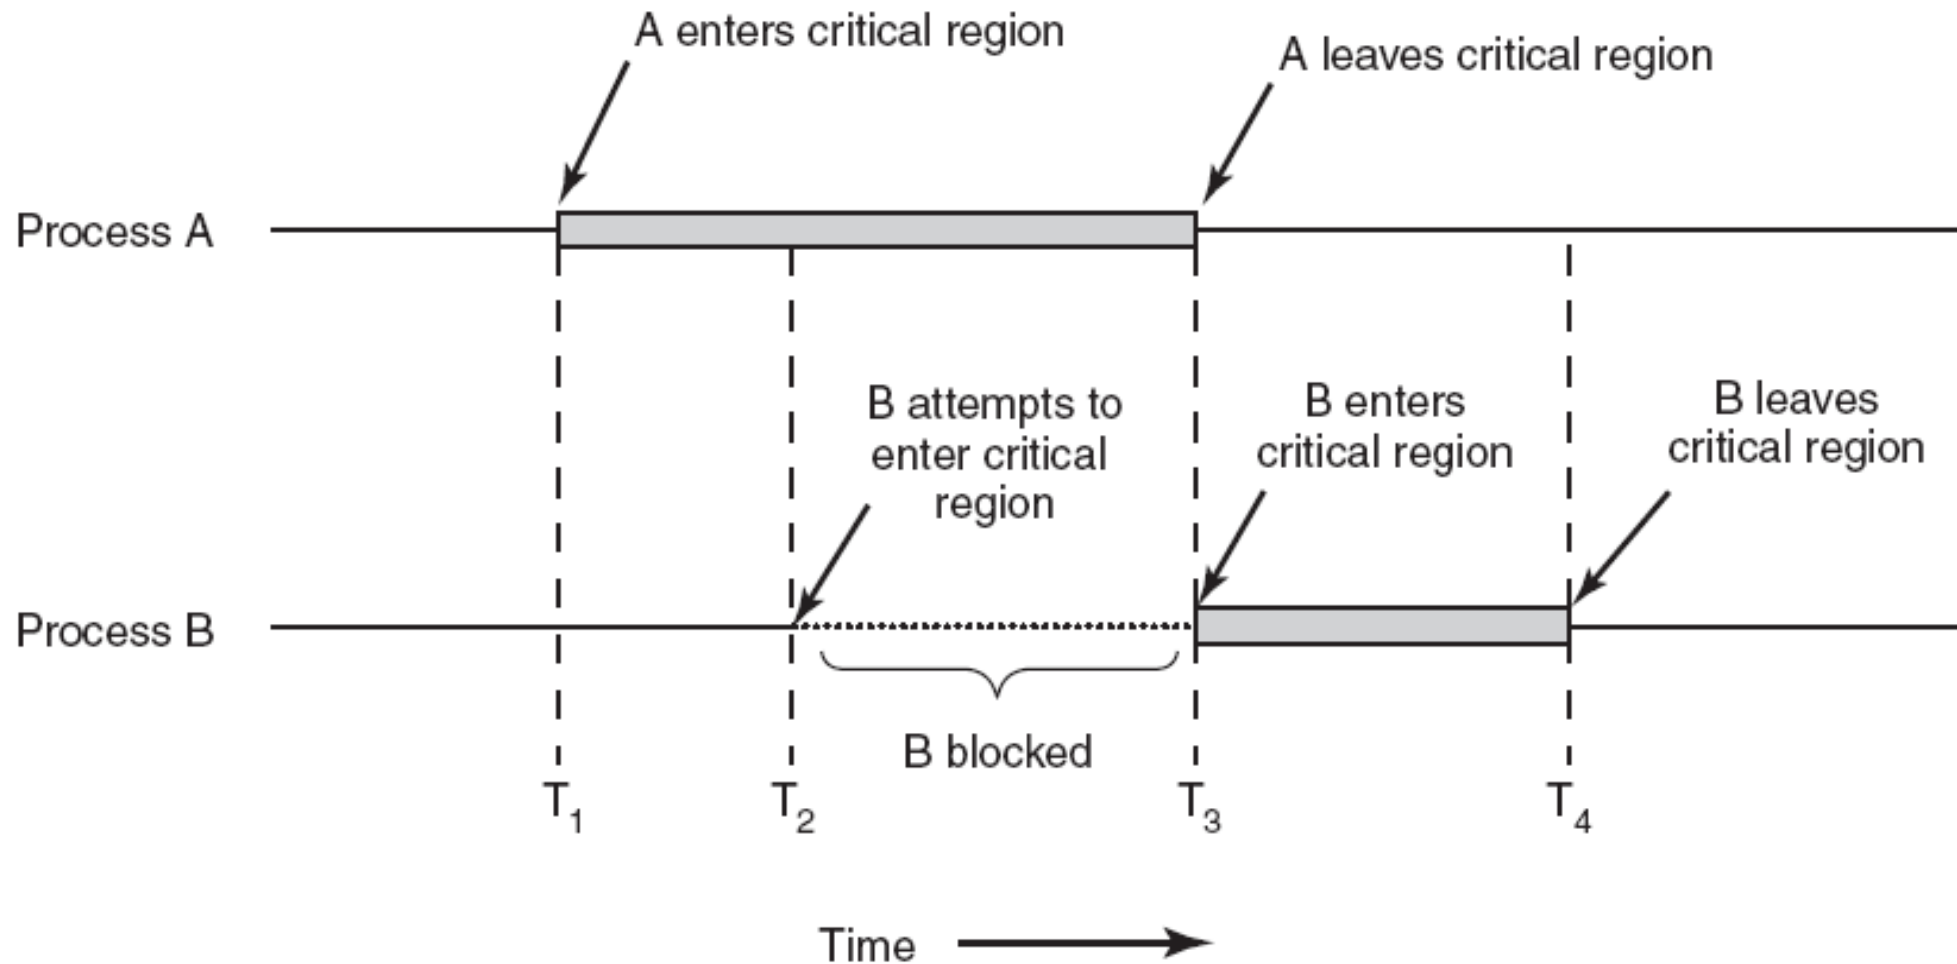
## Definition: Indivisible Operation /Atomic



An operation on a set of data items that cannot be executed concurrently, either with itself or with any other operation on a data item included in the set.

# Critical Sections (Regions)

- Mutual Exclusion & Critical Regions



# Synchronization Approaches

- Busy Waiting versus Blocking
- Hardware Support for Process Synchronization
- Algorithmic Approaches, Synchronization Primitives, and Concurrent Programming Constructs

# Busy Waiting

- A form of **Synchronization** in which a process repeatedly checks a condition until it becomes TRUE

```
While ( ' state FALSE' ) {  
    ;  
};
```

```
AA: cmp B, C  
    bne AA
```

while (some process is in a critical section on  $X_{\text{shared}}$   
or is executing an indivisible operation on  $X_{\text{shared}}$  )  
{ do nothing }


Critical Section or  
Indivisible Operation  
using  $X_{\text{shared}}$

- Can implement busy-waiting using **machine instructions**
- Inefficient in a multi-tasking system, but may be OK in multi-processor systems (1 process per processor).

# Strict Alternation or Hand-off Synchronization

- A **special case proposed solution** to the critical section problem.

(a) Process 0 executes when 'turn value 0'. (b) Process 1 executes when 'turn value 1'. In both cases, be sure to note the semicolons terminating the while statements



# Safe Access to a Shared Critical Sections

## Thread 1

CSEntry

Critical Section

CSExit

non-critical section

## Thread 2

CSEntry

Critical Section

CSExit

non-critical section

CSEntry represents the “protocol” required to safely enter the CS  
CSExit represents the “protocol” to safely release the CS

# Critical Section Control

## Using Simple Lock Variables

Boolean lock = 0; //unlocked, open

```
While (TRUE){  
    acquire (lock);  
    Critical_section();  
    release(lock);  
    Non-critical_section();  
}
```

```
While (TRUE){  
    acquire (lock);  
    Critical_section();  
    release(lock);  
    Non-critical_section();  
}
```

Can this implementation produce a safe solution?

# Implementing LOCKS: w/ Load+Store Code

uses a single shared lock variable

Lock Variable

0	1
Open	Closed
Unlocked	Locked
Free	Held

**Acquire** (perform Lock) : spin on a lock variable until it is unset, then set it in order to acquire lock

**Release** (Perform unlock): unset lock variable

```
boolean lock = false; // shared variable
```

```
void acquire(Boolean *lock) {  
    while (*lock) /* wait */ ;  
    *lock = true; }
```



CS

```
void release(Boolean *lock) {  
    *lock = false; }
```

Can this be used for Critical  
Section Control?



# A failed lock implementation

- Thread 1 spins, lock is released, ends spin
- Thread 1 interrupted just before setting flag
- Race condition has moved to the lock acquisition code!

Both threads grab the lock

Problem: Testing lock and setting lock are not atomic

Thread 1

```
acquire(*lock)
while (*lock)
```

Interrupt: Switch to Thread 2

Thread 2

```
acquire(*lock)
while (*lock);
*lock = 1;
```

Interrupt: switch to Thread 1

```
*lock = 1: //set flag to 1
```

# Hardware Support for Process Synchronization

- Need Indivisible instructions for the “wait implementation” (the CSEntry)
  - Avoid race conditions on memory locations (i.e., wait until safe)
- Used with a “lock variable” to implement indivisible operations and Critical Section (CS) control

Need to Know:

If CS access is  
open, then we take  
it and no one else  
sees it open


entry\_test:

```
If lock = closed  
    then goto entry_test;  
lock = closed;
```

{ Critical section or indivisible operation}

lock = open;

If OPEN, move forward with  
lock CLOSED for others



- *CS control and entry\_test* performed with an indivisible instruction
  - Test-and-set (TS) instruction
  - Swap instruction

# Synchronization Hardware

Many systems provide hardware support for critical section code and control

- **Disable Interrupts**

- Uniprocessors

- Currently running code would execute without preemption
    - Generally too inefficient on uniprocessor systems (all other activity on hold)

- **Does not work for multiprocessor systems**

Operating systems using this method are not broadly scalable

- Modern machines provide special atomic hardware instructions

- Atomic = non-interruptable

- Either test memory word AND set value
      - Or swap contents of two memory words
- Within 1 atomic instruction

Two Operations in  
One Instruction

# Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
} while (TRUE);
```

remainder section

A lock is just a variable whose value (state) is either:

Available / free /  
unlocked

or

Acquired / held /  
unlocked

Call lock() in order to **acquire a lock**. If the lock is free, the function returns. If not free, the caller waits until the lock is free and then acquires it.

Call unlock() to **release or free the lock** so that it is available for acquisition.

lock() == acquireLock()  
unlock() == releaseLock

# Building A Lock

## Goals of a lock implementation

- Mutual exclusion (obviously!)
- Fairness: all threads should eventually get the lock, and no thread should starve
- Low overhead: acquiring, releasing, and waiting for lock should not consume too many resources
- Implementation of locks are needed for both user space programs (e.g., pthreads library) and kernel code
- Implementing locks needs support from hardware and OS

# Solution: Hardware atomic instructions

- Very hard to ensure atomicity only in software
- Modern architectures provide hardware atomic instructions
- Example of an atomic instruction: test-and-set
  - Update a variable and return old value, all in one hardware instruction

```
int TestAndSet(int *old_ptr, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = new;      // store 'new' into old_ptr  
    return old;          // return the old value  
}
```

This is the operation of the  
hardware instruction TestAndSet

# The Test & Set 'Function'

TS() TestAndSet; where lock: Global, cc: Local

**TS(lock, cc):** <cc = lock; lock = TRUE>

"Atomic" instruction <...> indivisible operations

Global int lock = FALSE;

**CSEntry:**

```
int cc = FALSE;
TS(lock, cc);
While (cc){
    TS(lock, cc);
}
```

**SPIN-LOCK**

**CSExit:**

Lock = FALSE;

# The Test & Set Lock Instruction

- Spin-Lock Implementation with a TSL instruction

Note:  
REGISTER is  
“local”

enter\_region:

```
TSL REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

SPIN-LOCK

| copy lock to register and set lock to 1  
| was lock zero?  
| if it was nonzero, lock was set, so loop  
| return to caller; critical region entered

leave\_region:

```
MOVE LOCK,#0
RET
```

| store a 0 in lock  
| return to caller

TSL causes the CPU executing the instruction to lock the memory bus to prohibit other CPUs from accessing memory until complete. Ensures reads and writes are done atomically

This safe locking result would NOT be Accomplished by  
Disabling Interrupts in a multiprocessor system